Avneesh Jadhav

AI/ML Engineer

Email: avneeshjadhav1@gmail.com
Phone: +91 9545457385
Location: Pune, Maharashtra, India
LinkedIn: linkedin.com/in/avneeshjadhav/

SUMMARY

Working as an AI Engineer and Full-Stack Developer. Building production-grade AI systems and scalable full-stack solutions using Machine Learning, Generative AI, backend systems, and modern web stacks to turn complex ideas into practical solutions.

EXPERIENCE

AI Engineer and Full-Stack Developer | Kudzu Infotech (Sep 2026 - Present)

- Building and shipping AI-native applications and automation workflows
- Developing automation pipelines optimized for cloud deployments
- Improving performance, reliability, and deployment readiness for production AI systems

SKILLS

- **Programming Languages:** Python, C++, Rust, JavaScript, TypeScript
- **AI/ML:** Machine Learning, Deep Learning, Artificial Intelligence, Generative AI, Natural Language Processing, LLMs, LangChain, LangGraph, Agentic AI, Retrieval-Augmented Generation, Vector Databases, Prompt Engineering
- **Infrastructure and Systems:** Linux, Git, DevOps, CI/CD, Cloud Computing, APIs, Docker
- **Core:** Data Structures, Algorithms, Workflow Automation

EDUCATION

Bachelor of Engineering - Computer Engineering *(Nov 2022 - Jul 2026)*
University of Pune

PROJECTS

- **Agentic RAG System:** Developed a RAG pipeline with agentic workflow orchestration using Python, LangChain, LangGraph, and ChromaDB. Evaluated it using DeepEval framework with average scores: ~94.3% answer relevancy, ~98.1% faithfulness, ~100% context precision, ~100% context recall, and ~100% overall pass rate *(Jul 2026)* ***[GitHub]***
- **LLM from Scratch:** Built a 124 million parameter LLM transformer from scratch in PyTorch; trained it on ~2 billion tokens; achieved 14.8 validation perplexity and deployed it for inference via a serverless instance *(May 2026)* ***[GitHub]***
- **Project Vulcan:** Created a self-hostable AI platform with terminal integration, MCP tools, and browser control to automate daily workflows. Built using TypeScript, React, Rust, Axum, SQLite, and Docker *(Apr 2026)* ***[GitHub]***
- **LLM Inference Deployment:** Deployed Qwen3.8-27B on 2x Nvidia RTX 3090 GPUs using vLLM with tensor parallelism and FP8 quantization; built a benchmark suite and achieved ~38ms decode latency, ~8.4x faster prefill on prefix-cache hits, and ~$4/1M output tokens *(Nov 2026)* ***[GitHub]***

CERTIFICATIONS

- Machine Learning Specialization - Stanford University
- Fundamentals of Deep Learning - Nvidia
- Advanced Linux: The Linux Kernel - LinkedIn Learning
- Oracle AI Vector Search Professional - Oracle
- Claude Code in Action - Anthropic
- Model Context Protocol: Advanced Topics - Anthropic
- Advanced Rust: Managing Projects - LinkedIn Learning
- AMCAT Certified Software Engineer - AMCAT

ACHIEVEMENTS

- Research paper selected for the 11th ICTIS International Conference in Bangkok, Thailand
- Received golden badges on HackerRank for Python, C++, Java, and Problem Solving skills
- Successfully solved 270+ core DSA problems on LeetCode

Portfolio Website: avneeshjadhav.onrender.com/
GitHub: github.com/avneeshjadhav04